

```
In [ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: #upload the dataset on the left side in the folder for google collab first collected school data - http  
s://docs.google.com/spreadsheets/d/1nnxghnlQG2ZtCIplWtgLhrakPdrW0Uem5rewHSeRN2c/edit?usp=sharing  
  
import pandas as pd  
data = pd.read_csv("/content/school.csv")
```

```
In [ ]: data.drop("Index", axis=1, inplace=True)
```

In []: data

Out []:

	School Name	Total Enrolments	Student Teacher Ratio	One AP Test	Math Scoring	English Scoring	Graduation Rate	Minority Enrolment	Free Lunch Program	College Readiness
0	AAEC: Paradise Valley	322.0	NaN	NaN	42.0	NaN	97	28.6	NaN	NaN
1	AAEC: South Mountain	499.0	NaN	NaN	12.0	NaN	100	93.8	18.0	NaN
2	Accelerated :Learning Center	151.0	NaN	NaN	NaN	NaN	62	25.8	NaN	NaN
3	Alhambra High School	2395.0	18:1	5.0	21.0	NaN	70	95.8	78.0	14.5
4	Arcadia High School	1625.0	23:1	21.0	29.0	NaN	94	49.0	12.0	29.6
...
110	Roadrunner School	16.0	4:1	NaN	30.0	30	NaN	52.0	25.0	NaN
111	Ridgeline Academy Inc.	138.0	NaN	NaN	30.0	30	100	29.0	5.0	NaN
112	Pxu Digital Academy	1232.0	79:1	NaN	30.0	30	29	95.0	NaN	NaN
113	Pvonline	247.0	20:1	NaN	37.0	47	82	50.0	NaN	NaN
114	Premier Charter High School	146.0	NaN	NaN	30.0	30%	35	99.0	39.0	NaN

115 rows × 10 columns

```
In [ ]: missing_values_per_column = data.isnull().sum()

print("Number of missing values in each column:")
print(missing_values_per_column)
```

```
Number of missing values in each column:
School Name          0
Total Enrolments     1
Student Teacher Ratio 71
One AP Test          69
Math Scoring         38
English Scoring      54
Graduation Rate      10
Minority Enrolment   1
Free Lunch Program   35
College Readiness    70
dtype: int64
```

```
In [ ]: columns_to_check = ['School Name',      'Total Enrolments',      'Student Teacher Ratio',      'One AP Te
st',      'Math Scoring', 'English Scoring',      'Graduation Rate',      'Minority Enrolment',      'Free Lunc
h Program',      'College Readiness']

for col in columns_to_check:
    if data[col].astype(str).str.contains('%').any():
        print(f"Column '{col}' contains '%' signs.")
    else:
        print(f"{col} column has non numeric characters.")
```

```
School Name column has non numeric characters.
Total Enrolments column has non numeric characters.
Student Teacher Ratio column has non numeric characters.
One AP Test column has non numeric characters.
Math Scoring column has non numeric characters.
Column 'English Scoring' contains '%' signs.
Column 'Graduation Rate' contains '%' signs.
Minority Enrolment column has non numeric characters.
Free Lunch Program column has non numeric characters.
College Readiness column has non numeric characters.
```

```
In [ ]: from sklearn.preprocessing import MinMaxScaler

first_transform_df = data.copy()

first_transform_df['Graduation Rate'] = first_transform_df['Graduation Rate'].replace('%', '', regex=True).astype(float)
first_transform_df['English Scoring'] = first_transform_df['English Scoring'].replace('%', '', regex=True).astype(float)

numerical_columns = ['Total Enrolments', 'One AP Test', 'Math Scoring', 'English Scoring', 'Graduation Rate', 'Minority Enrolment', 'Free Lunch Program', 'College Readiness']

scaler = MinMaxScaler()
first_transform_df[numerical_columns] = scaler.fit_transform(first_transform_df[numerical_columns])
```

```
In [ ]: first_transform_df
```

```
Out[ ]:
```

	School Name	Total Enrolments	Student Teacher Ratio	One AP Test	Math Scoring	English Scoring	Graduation Rate	Minority Enrolment	Free Lunch Program	College Readiness
0	AAEC: Paradise Valley	0.106201	NaN	NaN	0.428571	NaN	0.97	0.286	NaN	NaN
1	AAEC: South Mountain	0.164578	NaN	NaN	0.122449	NaN	1.00	0.938	0.181818	NaN
2	Accelerated Learning Center	0.049802	NaN	NaN	NaN	NaN	0.62	0.258	NaN	NaN
3	Alhambra High School	0.789908	18:1	0.05	0.214286	NaN	0.70	0.958	0.787879	0.145
4	Arcadia High School	0.535950	23:1	0.21	0.295918	NaN	0.94	0.490	0.121212	0.296
...
110	Roadrunner School	0.005277	4:1	NaN	0.306122	0.38961	NaN	0.520	0.252525	NaN
111	Ridgeline Academy Inc.	0.045515	NaN	NaN	0.306122	0.38961	1.00	0.290	0.050505	NaN
112	Pxu Digital Academy	0.406332	79:1	NaN	0.306122	0.38961	0.29	0.950	NaN	NaN
113	Pvonline	0.081464	20:1	NaN	0.377551	0.61039	0.82	0.500	NaN	NaN
114	Premier Charter High School	0.048153	NaN	NaN	0.306122	0.38961	0.35	0.990	0.393939	NaN

115 rows × 10 columns

```
In [ ]: columns_to_drop = ['Student Teacher Ratio']
```

```
second_df = first_transform_df.drop(columns=columns_to_drop)
```

In []: second_df

Out []:

	School Name	Total Enrolments	One AP Test	Math Scoring	English Scoring	Graduation Rate	Minority Enrolment	Free Lunch Program	College Readiness
0	AAEC: Paradise Valley	0.106201	NaN	0.428571	NaN	0.97	0.286	NaN	NaN
1	AAEC: South Mountain	0.164578	NaN	0.122449	NaN	1.00	0.938	0.181818	NaN
2	Accelerated Learning Center	0.049802	NaN	NaN	NaN	0.62	0.258	NaN	NaN
3	Alhambra High School	0.789908	0.05	0.214286	NaN	0.70	0.958	0.787879	0.145
4	Arcadia High School	0.535950	0.21	0.295918	NaN	0.94	0.490	0.121212	0.296
...
110	Roadrunner School	0.005277	NaN	0.306122	0.38961	NaN	0.520	0.252525	NaN
111	Ridgeline Academy Inc.	0.045515	NaN	0.306122	0.38961	1.00	0.290	0.050505	NaN
112	Pxu Digital Academy	0.406332	NaN	0.306122	0.38961	0.29	0.950	NaN	NaN
113	Pvonline	0.081464	NaN	0.377551	0.61039	0.82	0.500	NaN	NaN
114	Premier Charter High School	0.048153	NaN	0.306122	0.38961	0.35	0.990	0.393939	NaN

115 rows × 9 columns

```
In [ ]: duplicate = second_df[second_df.duplicated(subset='School Name', keep=False)]
print("Duplicate Values:")
duplicate
```

Duplicate Values:

Out []:

	School Name	Total Enrolments	One AP Test	Math Scoring	English Scoring	Graduation Rate	Minority Enrolment	Free Lunch Program	College Readiness
83	Premier Charter High School	0.048153	0.0	0.000000	0.00000	0.35	0.993	0.393939	0.0
114	Premier Charter High School	0.048153	NaN	0.306122	0.38961	0.35	0.990	0.393939	NaN

```
In [ ]: third_df = second_df.groupby('School Name').agg(lambda x: x.max()).reset_index()
```

```
In [ ]: duplicate = third_df[third_df.duplicated(subset='School Name', keep=False)]
print("Duplicate Values:")
duplicate
```

Duplicate Values:

Out []:

	School Name	Total Enrolments	One AP Test	Math Scoring	English Scoring	Graduation Rate	Minority Enrolment	Free Lunch Program	College Readiness
--	-------------	------------------	-------------	--------------	-----------------	-----------------	--------------------	--------------------	-------------------

```
In [ ]: mean_columns = ['One AP Test', 'English Scoring', 'College Readiness']
for column in mean_columns:
    mean_value = third_df[column].mean()
    third_df[column].fillna(mean_value, inplace=True)
```

In []: third_df

Out []:

	School Name	Total Enrolments	One AP Test	Math Scoring	English Scoring	Graduation Rate	Minority Enrolment	Free Lunch Program	College Readiness
0	AAEC: Paradise Valley	0.106201	0.251522	0.428571	0.391991	0.97	0.286	NaN	0.219622
1	AAEC: South Mountain	0.164578	0.251522	0.122449	0.391991	1.00	0.938	0.181818	0.219622
2	ASU Preparatory Academy-Phoenix High School	0.108509	0.040000	0.275510	0.391991	0.97	0.933	0.181818	0.074000
3	Accelarated :Learning Center	0.049802	0.251522	NaN	0.391991	0.62	0.258	NaN	0.219622
4	Alhambra High School	0.789908	0.050000	0.214286	0.391991	0.70	0.958	0.787879	0.145000
...
109	West-Mec - Paradise Valley High School	0.008245	0.251522	0.306122	0.389610	0.00	0.640	0.000000	0.219622
110	West-Mec - Pinnacle High School	0.008575	0.251522	0.306122	0.389610	0.00	0.346	0.000000	0.219622
111	Western School of Science and Technology	0.114116	0.430000	0.173469	0.415584	0.95	0.991	0.262626	0.135000
112	Westland School	0.023747	0.251522	0.306122	0.389610	0.50	0.831	0.767677	0.219622
113	e-Institute at Metro	0.019129	0.251522	NaN	0.391991	0.05	0.931	NaN	0.219622

114 rows × 9 columns


```
In [ ]: third_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 114 entries, 0 to 113
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   School Name           114 non-null   object
1   Total Enrolments      113 non-null   float64
2   One AP Test           114 non-null   float64
3   Math Scoring          76 non-null    float64
4   English Scoring       114 non-null   float64
5   Graduation Rate       104 non-null   float64
6   Minority Enrolment    113 non-null   float64
7   Free Lunch Program    79 non-null    float64
8   College Readiness     114 non-null   float64
dtypes: float64(8), object(1)
memory usage: 8.1+ KB
```

```
In [ ]: cols_to_drop_na = ['Total Enrolments', 'Math Scoring', 'Graduation Rate', 'Minority Enrolment', 'Free Lunch Program']
third_df.dropna(subset=cols_to_drop_na, inplace=True)
```

```
In [ ]: print("Dimensions of the DataFrame after being cleaned and transformed as per the Plan: {}".format(third_df.shape))
```

Dimensions of the DataFrame after being cleaned and transformed as per the Plan: (62, 9)

```
In [ ]: final_df = third_df
```

```
In [ ]: final_df.to_csv('final_df_extra_credit.csv', index=False)
```

```
In [ ]: from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LinearRegression
        from sklearn.metrics import mean_squared_error, r2_score

        X = final_df[['Total Enrolments', 'One AP Test', 'English Scoring', 'Graduation Rate', 'Minority Enrolment', 'Free Lunch Program', 'College Readiness']]
        y = final_df['Math Scoring']

        model = LinearRegression()

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)

        model.fit(X_train, y_train)

        y_pred = model.predict(X_test)
```

```
In [ ]: mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
print('Mean Squared Error (MSE) – Measures the average of the squares of the errors')
print(mse, '\n')
print('Root Mean Squared Error (RMSE)– Represents the square root of the MSE')
print(rmse, '\n')
print('R-squared (R²) – Indicates the proportion of variance in the dependent variable that is predictable from the independent variables')
print(r2, '\n')

print("Coefficient Interpretation:")
for feature, coefficient in zip(X.columns, model.coef_):
    if coefficient > 0:
        print(f"A one-unit increase in {feature} is associated with an increase of in Math Scoring by – {coefficient:.4f}")
    elif coefficient < 0:
        print(f"A one-unit increase in {feature} is associated with a decrease of in Math Scoring by – {abs(coefficient):.4f}")
    else:
        print(f"The {feature} does not have a significant impact on Math Scoring.")

print('\nIntercept (baseline Math Scoring when all other factors are zero):', model.intercept_)
```

Mean Squared Error (MSE) – Measures the average of the squares of the errors
0.027787678827900356

Root Mean Squared Error (RMSE)– Represents the square root of the MSE
0.1666963671706746

R-squared (R^2) – Indicates the proportion of variance in the dependent variable that is predictable from the independent variables
0.09742671988922114

Coefficient Interpretation:

A one-unit increase in Total Enrolments is associated with a decrease of in Math Scoring by – 0.0586

A one-unit increase in One AP Test is associated with a decrease of in Math Scoring by – 0.0260

A one-unit increase in English Scoring is associated with an increase of in Math Scoring by – 0.6945

A one-unit increase in Graduation Rate is associated with a decrease of in Math Scoring by – 0.0412

A one-unit increase in Minority Enrolment is associated with a decrease of in Math Scoring by – 0.0545

A one-unit increase in Free Lunch Program is associated with an increase of in Math Scoring by – 0.0170

A one-unit increase in College Readiness is associated with an increase of in Math Scoring by – 0.1118

Intercept (baseline Math Scoring when all other factors are zero): 0.043759542802430995

```
In [ ]: from sklearn.ensemble import BaggingRegressor

base_model = LinearRegression()

bagging_model = BaggingRegressor(estimator=base_model, n_estimators=10, random_state=42)

bagging_model.fit(X_train, y_train)

y_pred = bagging_model.predict(X_test)
```

```
In [ ]: mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
print('Mean Squared Error (MSE) – Measures the average of the squares of the errors')
print(mse, '\n')
print('Root Mean Squared Error (RMSE)– Represents the square root of the MSE')
print(rmse, '\n')
print('R-squared (R²) – Indicates the proportion of variance in the dependent variable that is predictable from the independent variables')
print(r2, '\n')

print("Coefficient Interpretation:")
for feature, coefficient in zip(X.columns, model.coef_):
    if coefficient > 0:
        print(f"A one-unit increase in {feature} is associated with an increase of in Math Scoring by – {coefficient:.4f} increase")
    elif coefficient < 0:
        print(f"A one-unit increase in {feature} is associated with a decrease of in Math Scoring by – {abs(coefficient):.4f} decrease")
    else:
        print(f"The {feature} does not have a significant impact on Math Scoring.")

print('\nIntercept (baseline Math Scoring when all other factors are zero):', model.intercept_)
```

Mean Squared Error (MSE) – Measures the average of the squares of the errors
0.027787678827900356

Root Mean Squared Error (RMSE)– Represents the square root of the MSE
0.1666963671706746

R-squared (R^2) – Indicates the proportion of variance in the dependent variable that is predictable from the independent variables
0.09742671988922114

Coefficient Interpretation:

A one-unit increase in Total Enrolments is associated with a decrease of in Math Scoring by – 0.0586 decrease

A one-unit increase in One AP Test is associated with a decrease of in Math Scoring by – 0.0260 decrease

A one-unit increase in English Scoring is associated with an increase of in Math Scoring by – 0.6945 increase

A one-unit increase in Graduation Rate is associated with a decrease of in Math Scoring by – 0.0412 decrease

A one-unit increase in Minority Enrolment is associated with a decrease of in Math Scoring by – 0.0545 decrease

A one-unit increase in Free Lunch Program is associated with an increase of in Math Scoring by – 0.0170 increase

A one-unit increase in College Readiness is associated with an increase of in Math Scoring by – 0.1118 increase

Intercept (baseline Math Scoring when all other factors are zero): 0.043759542802430995

```
In [ ]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, c='blue', alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2)
plt.xlabel('Actual Math Scores')
plt.ylabel('Predicted Math Scores')
plt.title('Actual vs Predicted Math Scores')
plt.grid(True)
plt.show()

import matplotlib.pyplot as plt

sorted_indices = np.argsort(y_test)
y_test_sorted = y_test.iloc[sorted_indices]
y_pred_sorted = y_pred[sorted_indices]

plt.figure(figsize=(10, 6))
plt.plot(range(len(y_test)), y_test_sorted, marker='o', color='blue', label='Actual', linewidth=2)
plt.plot(range(len(y_test)), y_pred_sorted, marker='x', color='red', label='Predicted', linewidth=2)
plt.xlabel('Instance')
plt.ylabel('Math Scores')
plt.title('Actual vs Predicted Math Scores')
plt.legend()
plt.grid(True)
plt.show()
```



